

Birds of a feather flock together

Clusters



Susan Holmes Wolfgang Huber Laura Symul October 9, 2019

Clustering =

Finding a latent or hidden variable which is not necessarily provided/accessible

Clusters =

Groups in the population (of individuals, cells, gene expression profiles) in which the **individual elements are similar to each others** and **different from the elements of other groups**.



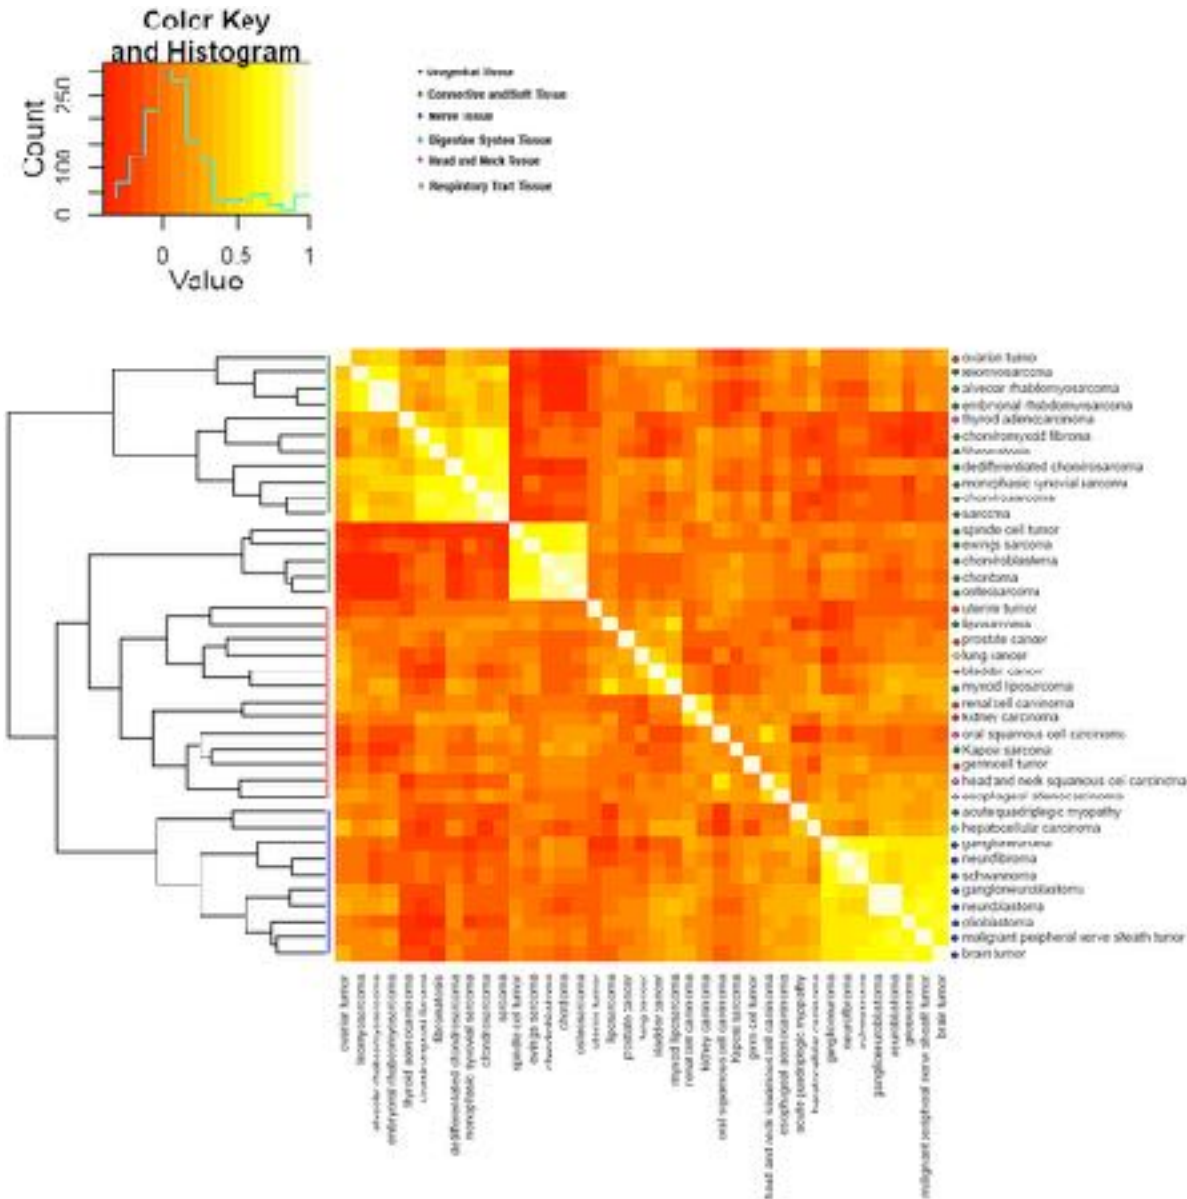
Some **measurements**
(imaging, single cell RNA-seq, etc.)

Groups = cell types


What is the central statistical concept behind any clustering?

Metric of (dis)similarity : a distance

Clusters frequently occur in biology.
Cell type, tissue, tumor type, ethnicities, etc.



Data Augmentation is the aim of clustering



data augmentation methods are methods that add useful but unknown components to the data

Goals for this lecture


1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as *k-means* or *k-medoids*
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called *hierarchical clustering*.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.

How (dis)similar are 2 elements?



(dis)similarity measure = **distance**

How (dis)similar are 2 elements?



How (dis)similar are 2 elements?

0	1	0	0	0	1	0	0	1	0	0	1
0	0	1	0	0	1	0	0	1	0	0	1
<hr/>											
0	0	1	0	0	1	0	0	1	0	0	1
1	1	0	1	1	0	1	1	0	1	1	0

The choice or definition of distance depends on the data

Euclidean distance (L2)

$$d(A, B) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_p - b_p)^2}.$$

Manhattan distance (L1)

$$d(A, B) = |a_1 - b_1| + |a_2 - b_2| + \dots + |a_p - b_p|.$$

Maximum distance (L ∞)

$$d_\infty(A, B) = \max_i |a_i - b_i|.$$

Minkowski distance (L m)

$$d(A, B) = ((a_1 - b_1)^m + (a_2 - b_2)^m + \dots + (a_p - b_p)^m)^{\frac{1}{m}}.$$

Edit (Hamming) distance

Binary distance

Jaccard distance

$$d_J(S, T) = 1 - J(S, T) = \frac{f_{01} + f_{10}}{f_{01} + f_{10} + f_{11}}.$$





Correlation-based distance

$$d(A, B) = \sqrt{2(1 - \text{cor}(A, B))}.$$


Weighted Euclidean distance

Mahalanobis distance

...




Euclidean distance



$$d(A, B) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_p - b_p)^2}.$$

Weighted Euclidean distance

Is the **red point**, a “new” datapoint, closer to the cluster on the right or the cluster on the left?



The weight on the dimensions depends on the cluster

Binary distance

0	1	0	0	0	1	0	0	1	0	0	1
0	0	1	0	0	1	0	0	1	0	0	1
<hr/>											
0	0	1	0	0	1	0	0	1	0	0	1
1	1	0	1	1	0	1	1	0	1	1	0

$$1- \frac{\text{The sum of elements that are the same}}{\text{The number of elements}}$$

Jaccard distance

0	1	0	0	0	1	0	0	1	0	0	1
0	0	1	0	0	1	0	0	1	0	0	1
<hr/>											
0	0	1	0	0	1	0	0	1	0	0	1
1	1	0	1	1	0	1	1	0	1	1	0

The **co-occurrence** is more important than the co-absence

$$1 - \frac{\text{The sum of 1's that overlap}}{\text{The union of 1's}}$$

Initially developed for ecological data for the occurrence of traits or features

Edit (Hamming) distance

$$1 \left(\begin{array}{l} \text{KAROLIN} \\ \text{KATHRIN} \\ \text{CAROLIN} \end{array} \right)^3$$

The number of edits that needs to be done.

This could be applied to nucleotide or amino acid sequences...

-	m	t	v	e	l	p	s	t	q	r	a	l	v	f	d	t	w	n	g	p	l	e	v	r	q	v	p	v	p	s	p	a	d	d	e	i	l
m	s	t	a	g	k	v	i	k	c	k	a	a	v	l	w	e	l	k	k	p	f	s	i	e	e	v	e	v	a	p	p	k	a	h	e	v	r
-	-	-	m	s	i	p	e	t	q	k	g	v	i	f	y	e	s	h	g	k	l	e	y	k	d	i	p	v	p	k	p	k	a	n	e	l	l
-	-	m	a	n	q	v	i	r	c	k	a	a	v	a	w	e	a	g	k	p	l	s	i	e	e	i	e	v	a	p	p	k	a	h	e	v	r
m	g	t	q	g	k	v	i	k	c	k	a	a	i	a	w	k	t	g	s	p	l	c	i	e	e	i	e	v	s	p	p	k	a	c	e	v	r

...although in that case, the **different character substitutions** are usually associated with **different contributions to the distance** (to account for physical or evolutionary similarity), and **deletions and insertions** may also be allowed

Distances in R

```
mx = c(0, 0, 0, 1, 1, 1)
my = c(1, 0, 1, 1, 0, 1)
mz = c(1, 1, 1, 0, 1, 1)
mat = rbind(mx, my, mz)
```

Euclidean

```
dist(mat)
```

```
##           mx           my
## my 1.732051
## mz 2.000000 1.732051
```

Binary

```
dist(mat, method = "binary")
```

```
##           mx           my
## my 0.6000000
## mz 0.6666667 0.5000000
```

```
mut = read.csv("../data/HIVmutations.csv")
mut[1:3, 10:16]
```

```
##   p32I p33F p34Q p35G p43T p46I p46L
## 1    0    1    0    0    0    0    0
## 2    0    1    0    0    0    1    0
## 3    0    1    0    0    0    0    0
```

```
library("vegan")
mutJ = vegdist(mut, "jaccard")
mutC = sqrt(2 * (1 - cor(t(mut))))
```

Jaccard

```
mutJ
```

```
##           1    2    3    4
## 2 0.800
## 3 0.750 0.889
## 4 0.900 0.778 0.846
## 5 1.000 0.800 0.889 0.900
```

Correlation-based

```
as.dist(mutC)
```

```
##           1    2    3    4
## 2 1.19
## 3 1.10 1.30
## 4 1.32 1.13 1.30
## 5 1.45 1.19 1.30 1.32
```

Distances in R

phyloseq package

phyloseq: analyze microbiome census data using R

The phyloseq package is a tool to import, store, analyze, and graphically display complex phylogenetic sequencing data that has already been clustered into Operational Taxonomic Units (OTUs), especially when there is associated sample data, phylogenetic tree, and/or taxonomic assignment of the OTUs.

```
> library("phyloseq")
> distanceMethodList
$UniFrac
[1] "unifrac" "wunifrac"

$DPCoA
[1] "dpcoa"

$JSD
[1] "jsd"

$svegdist
[1] "manhattan" "euclidean" "canberra" "bray" "kulczynski" "jaccard" "gower" "altGower" "morisita"
[10] "horn" "mountford" "raup" "binomial" "chao" "cao"

$betadiver
[1] "w" "-1" "c" "wb" "r" "l" "e" "t" "me" "j" "sor" "m" "-2" "co" "cc" "g" "-3" "l" "19" "hk"
[21] "rlb" "sim" "gl" "z"

$dist
[1] "maximum" "binary" "minkowski"

$designndist
[1] "ANY"

> |
```

Almost 50 different distances

If you **really** need to define a **new distance**, it must...

Definition [\[edit \]](#)

A **metric space** is an **ordered pair** (M, d) where M is a set and d is a **metric** on M , i.e., a **function**

$$d: M \times M \rightarrow \mathbb{R}$$

such that for any $x, y, z \in M$, the following holds:^[2]

1. $d(x, y) = 0 \Leftrightarrow x = y$ **identity of indiscernibles**
2. $d(x, y) = d(y, x)$ **symmetry**
3. $d(x, z) \leq d(x, y) + d(y, z)$ **subadditivity** or **triangle inequality**

Given the above three axioms, we also have that $d(x, y) \geq 0$ for any $x, y \in M$. This is deduced as follows:

https://en.wikipedia.org/wiki/Metric_space

Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as *k-means* or *k-medoids*
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called *hierarchical clustering*.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.

Non-parametric clustering methods

(k-methods, top-down)

k-medoids

PAM (Partitioning Around (k) Medoids)

k-means

1. Starts from a matrix of p features measured on a set of n observations.
2. **Randomly pick k distinct cluster centers** out of the n observations (“seeds”).
3. **Assign each of the remaining observations** to the group to whose center it is the closest.
4. **For each group, choose a new center** from the observations in the group, such that the sum of the distances of group members to the center is minimal; this is called the *medoid*.
5. **Repeat Steps 3 and 4 until the groups stabilize.**




Non-parametric clustering methods

(k-methods, top-down)

The difference between k-medoids and k-means is at step 4, when choosing the new centers

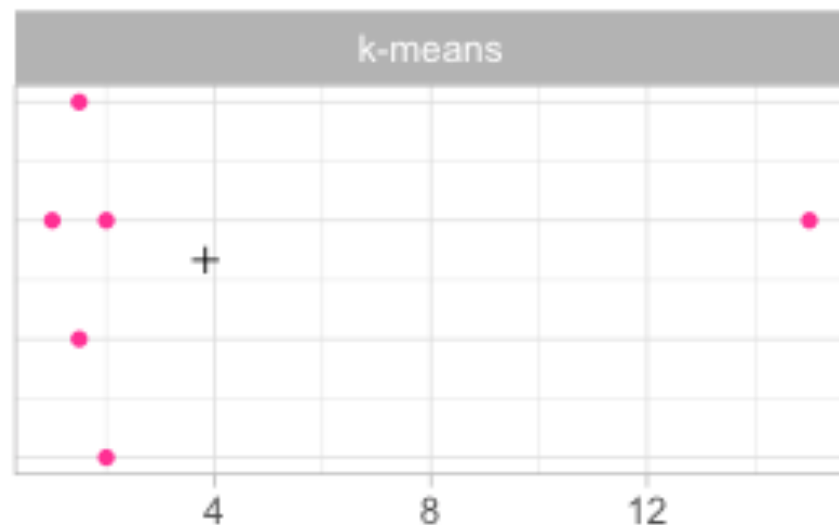
k-medoids

The new centers are selected **among the observations**



k-means

The new centers are **computed** as the **arithmetic mean** between the observations of each group



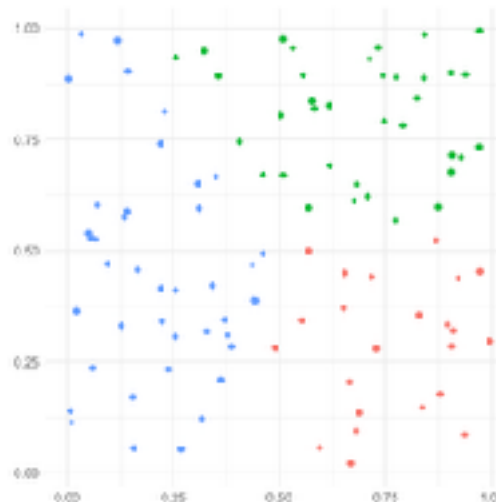

Non-parametric clustering methods (k-methods, top-down)

```
Xmat=matrix(runif(200),ncol=2)
nk=3
cents=matrix(runif(2*nk),ncol=2)

out=kmeans(Xmat,centers = nk) # give the # of clusters
out=kmeans(Xmat,centers = cents) # give the seeds (initial center of clusters)

X = data.frame(x = Xmat[,1], y = Xmat[,2], cluster = factor(out$cluster))

g = ggplot(X, aes(x = x, y = y)) + xlab("") + ylab("") + guides(col = FALSE)
g + geom_point()
g + geom_point(aes(col = cluster))
```



```
out1=kmeans(Xmat,cents,iter.max=1)
out2=kmeans(Xmat,cents,iter.max=2)
out3=kmeans(Xmat,cents,iter.max=3)
```

Non-parametric clustering methods

(k-methods, top-down)

1. Starts from a matrix of p features measured on a set of n observations.
2. **Randomly pick k distinct cluster centers** out of the n observations (“seeds”).
3. **Assign each of the remaining observations** to the group to whose center it is the closest.
4. **For each group, choose a new center** from the observations in the group, such that the sum of the distances of group members to the center is minimal; this is called the *medoid*.
5. **Repeat Steps 3 and 4 until the groups stabilize.**

Is the method robust?

i.e. does it give always the same result?

How can we improve that?



Strong Forms & Tight Clusters

Dynamical Clusters

Repeats the process many times with **different seeds** (initial centers) to build `**strong forms**' which are groups of observations that end up in the same classes for most possible initial configurations.

Diday and Brito, 1989

Resampling

Repeats the process many times on subsamples of the dataset to create `**tight clusters**' which are groups of observations that are almost always grouped together.

Tseng and Wong, 2015

Data preprocessing matters


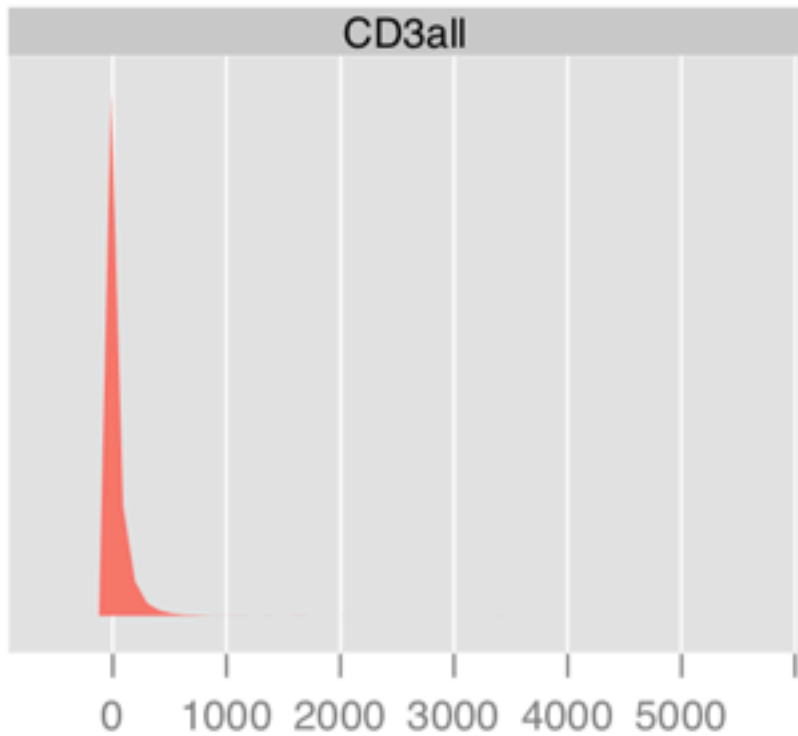



Figure 5.11: Cell measurements that show clear clustering in two dimensions.

Different scales


Raw




asinh



Density-based clustering



Density-based clustering



2 parameters:

Eps (ϵ) : the maximum distance between 2 points so that they are considered as “reachable” from one another


minPts: the minimum number of points for a cluster to be created.

It classifies each data-point into

- core points
- directly reachable points
- noise / outliers

Algorithm: **DBSCAN**

Density-based spatial clustering
of applications with noise



Martin Ester, Hans-Peter Kriegel, Jörg Sander and Xiaowei Xu, 1996

Package and function dbscan in R

<https://en.wikipedia.org/wiki/DBSCAN>

How does density-based clustering (dbscan) work ?

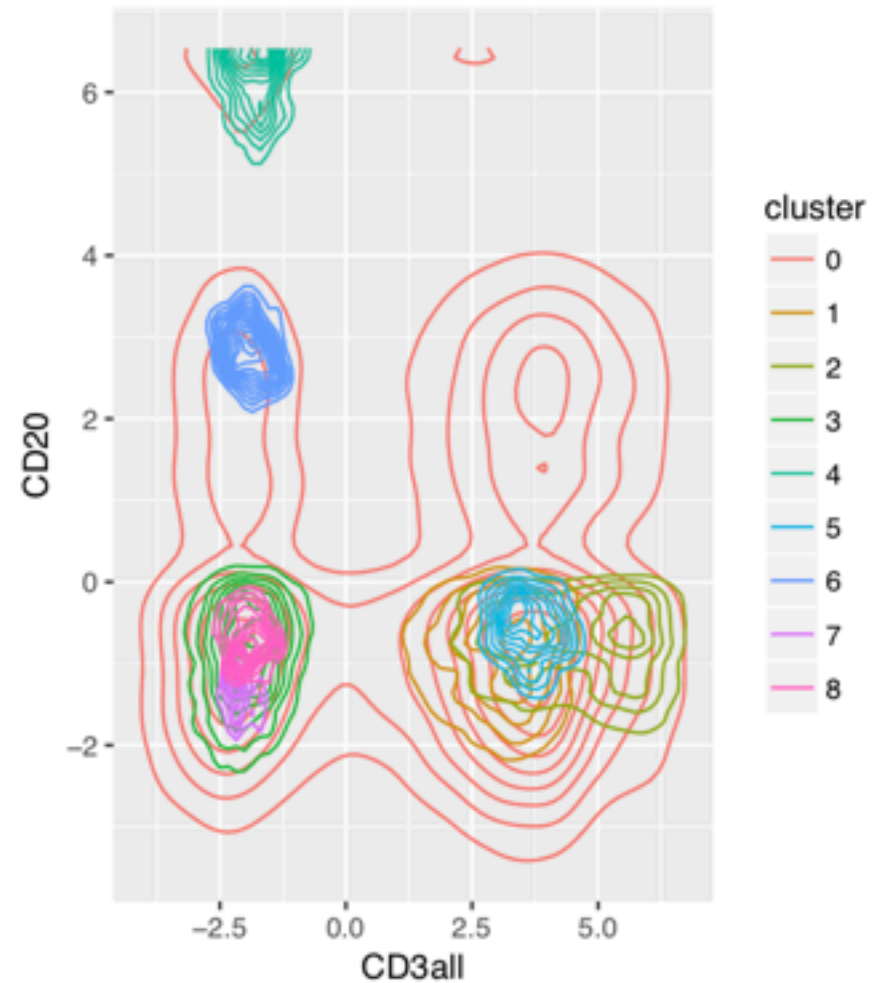
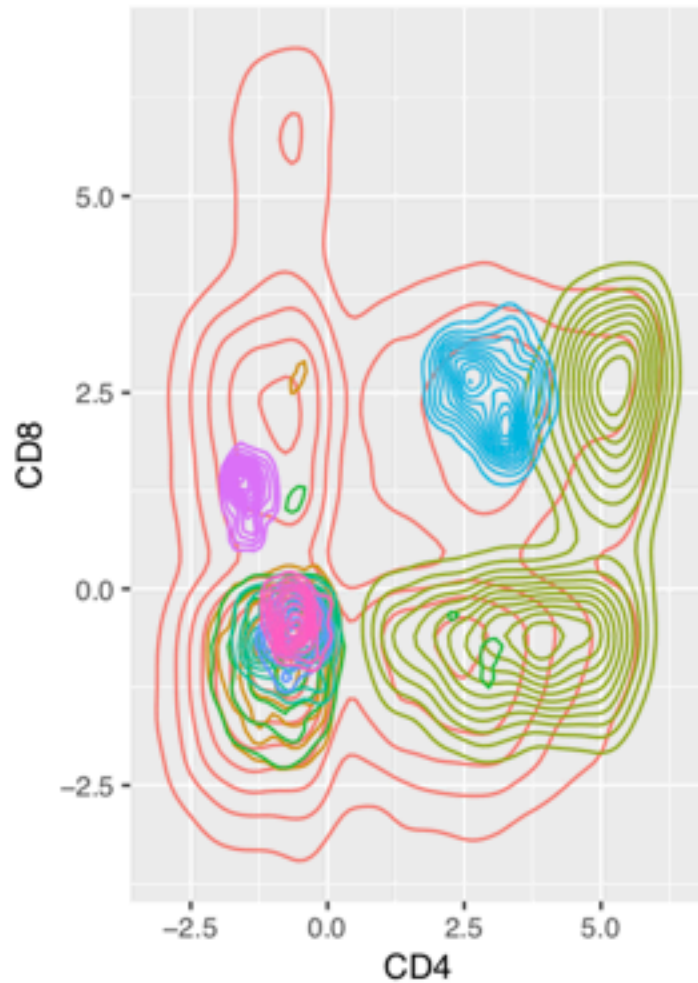
The dbscan method clusters points in dense regions according to the **density-connectedness** criterion. It looks at small neighborhood spheres of radius ϵ to see if points are connected.

The building block of dbscan is the concept of density-reachability: a point q is directly **density-reachable** from a point p if it is not further away than a given threshold ϵ , and if p is surrounded by sufficiently many points such that one may consider p (and q) be part of a dense region. We say that q is **density-reachable** from p if there is a sequence of points p_1, \dots, p_n with $p_1 = p$ and $p_n = q$, so that each p_{i+1} is directly density-reachable from p_i .

A **cluster** is then a subset of points that satisfy the following properties:

1. All points within the cluster are mutually density-connected.
2. If a point is density-connected to any point of the cluster, it is part of the cluster as well.
3. Groups of points must have at least `MinPts` points to count as a cluster.





Density-based clustering



Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as *k-means* or *k-medoids*
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called *hierarchical clustering*.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.

Hierarchical Clustering (bottom-up)




Hierarchical Clustering (bottom-up)

```
set.seed(1)
Xmat=matrix(runif(20),ncol=2)
d = dist(Xmat)
```

```
> round(d,3)
```



	1	2	3	4	5	6	7	8	9
2	0.111								
3	0.571	0.549							
4	0.667	0.575	0.452						
5	0.567	0.617	0.380	0.805					
6	0.697	0.617	0.377	0.114	0.748				
7	0.850	0.788	0.373	0.336	0.745	0.225			
8	0.880	0.865	0.317	0.656	0.510	0.548	0.395		
9	0.403	0.328	0.312	0.279	0.578	0.294	0.462	0.613	
10	0.607	0.676	0.519	0.933	0.140	0.882	0.885	0.636	0.693



Hierarchical Clustering (bottom-up)

```
set.seed(1)
Xmat=matrix(runif(20),ncol=2)
d = dist(Xmat)
```


```
clust = hclust(d, method = "ward.D")
plot(clust)
```



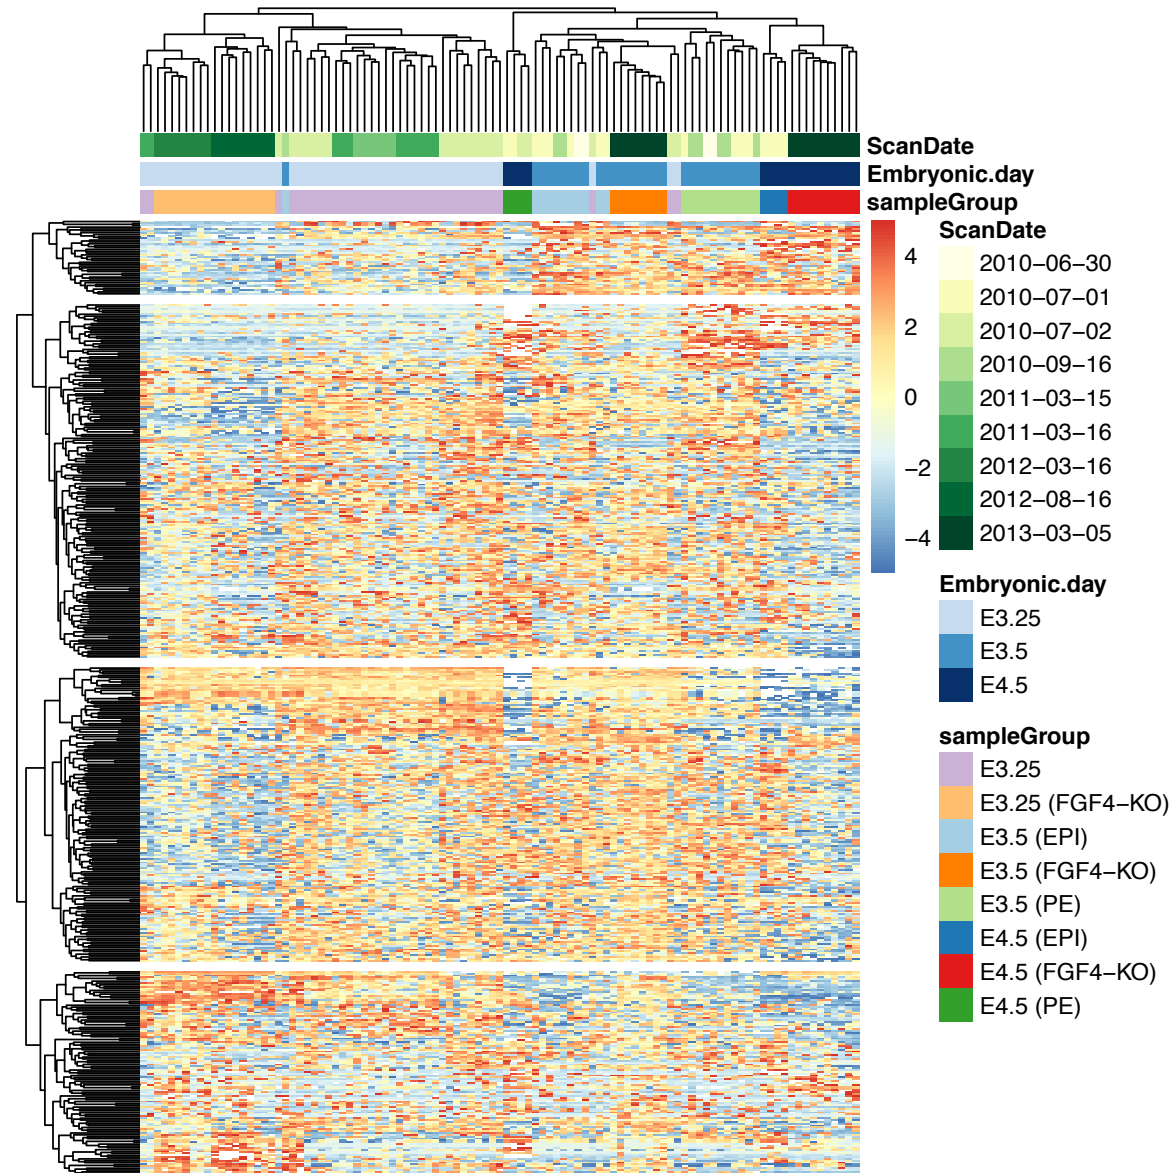
Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as ***k-means*** or ***k-medoids***
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called ***hierarchical clustering***.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.


The ordering of the branches is not unique



The order of dendrogram branches is not unique



The ordering of the branches is not unique




What matters is the “**order of the agglomerations**” and the **length between leaves**

Different options to agglomerate


```
clust = hclust(d, method = "ward.D")  
plot(clust)
```

Cluster Dendrogram



`hclust(d, "ward.D")`

Cluster Dendrogram




`hclust(d, "complete")`


Can you think of different ways to agglomerate a new point (or subgroup) ?

Different options to agglomerate


Single linkage




Good for recognizing the number of clusters
But “combs”




Maximal linkage




Compact classes
But one observation can alter groups




Medoids




More robust to outliers




Different options lead to **different tree shapes**





Single Linkage



Complete Linkage



Average Linkage




of clusters \leftrightarrow max dist between groups

You cut the tree such that

- you have a given number of cluster
- at a given height (distance between the groups)

How do you choose the number of clusters?



Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as ***k-means*** or ***k-medoids***
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called ***hierarchical clustering***.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.



Clustering methods **WILL** find clusters

... even if there are none

How do we validate the clusters?

How do we find the optimal number of cluster?

We want to

Minimize the distance between the points of a cluster (cohesion),

Maximizing the distance between clusters (separation).


There are different ways to assess this.

- > WSS (Within-group sum of squares)
- Elbow method
- Calinski-Harabasz index
- Gap statistic


- Silhouette method

WSS


$$WSS_k = \sum_{\ell=1}^k \sum_{x_i \in C_\ell} d^2(x_i, \bar{x}_\ell)$$




Data



2 clusters



3 clusters



Calinski-Harabasz index

We want to maximize


$$\text{CH}(k) = \frac{\text{Between-groups sum of squares}}{\text{Within-groups sum of squares}}$$

We want to maximize

We want to minimize

$$\text{CH}(k) = \frac{\text{BSS}_k}{\text{WSS}_k} \times \frac{N - k}{N - 1} \quad \text{where} \quad \text{BSS}_k = \sum_{\ell=1}^k n_{\ell} (\bar{x}_{\ell} - \bar{x})^2,$$

where \bar{x} is the overall center of mass (average point).



The Gap statistic

Algorithm for computing the gap statistic (Tibshirani, Walther, and Hastie 2001):

1. Cluster the data with k clusters and compute WSS_k for the various choices of k .
2. Generate B plausible reference data sets, using Monte Carlo sampling from a homogeneous distribution and redo Step 1 above for these new simulated data. This results in B new within-sum-of-squares for simulated data W_{kb}^* , for $b = 1, \dots, B$.
3. Compute the $gap(k)$ -statistic:

$$gap(k) = l_k - \log WSS_k \quad \text{with} \quad l_k = \frac{1}{B} \sum_{b=1}^B \log W_{kb}^*$$

Note that the first term is expected to be bigger than the second one if the clustering is good (i.e., the WSS is smaller); thus the gap statistic will be mostly positive and we are looking for its highest value.

4. We can use the standard deviation

$$sd_k^2 = \frac{1}{B-1} \sum_{b=1}^B (\log(W_{kb}^*) - l_k)^2$$


to help choose the best k . Several choices are available, for instance, to choose the smallest k such that

$$gap(k) \geq gap(k+1) - s'_{k+1} \quad \text{where} \quad s'_{k+1} = sd_{k+1} \sqrt{1 + 1/B}.$$

The packages **cluster** and **clusterCrit** provide implementations.



Monte Carlo



The Silhouette method

For each datapoint we compare


- the **average distance** to all the points of **its own cluster**
- the **average distance** to all the points of the **nearest cluster**

a_i = average of —

b_i = average of —

$$s_i = \frac{b_i - a_i}{\max\{a_i, b_i\}}$$

$$-1 \leq s_i \leq 1$$





Good score

Bad score:
Not very well integrated in the cluster

The Silhouette method



Silhouette analysis for KMeans clustering on sample data with $n_clusters = 3$



Silhouette analysis for KMeans clustering on sample data with $n_clusters = 4$




Silhouette analysis for KMeans clustering on sample data with $n_clusters = 5$



The bootstrap to validate clusters

How do we apply the bootstrap?





We create new datasets by subsampling the original dataset and evaluate how clusters are conserved.

Bootstrap for Clustering methods



Monte Carlo



Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as *k-means* or *k-medoids*
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called *hierarchical clustering*.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.

Some extra notes

There are already too many!

<https://cran.r-project.org/web/views/Cluster.html>

- Package [factanal](#) implements a likelihood-based factor analysis via the *factanal* function. This function also prints for each data point the factor loadings against the observed factor loadings with a higher loadings value. The cluster centroids of this factor loadings are obtained by using an iterative or automatic selection method.
- Package [factanal](#) provides alternative implementation of 4 means and approximate hierarchical clustering.
- Package [factanal](#) provides a new algorithm to reconstruct two-dimensional data.
- Package [factanal](#) implements clustering algorithms for business analysis like 'k-means' and 'mclust'.
- Package [factanal](#) clusters 3-dimensional data into four subgroups based on a categorical form of Chi and Hot's (2008) late learning method.
- Package [factanal](#) implements hierarchical clustering for both hierarchical and partitioning-based methods.
- Package [factanal](#) implements a cluster algorithm that is used on equal functions and hierarchy alone to group observations according to the multivariate dependence structure of the generating process without any assumptions on the margins.
- These cluster algorithms are available in package [factanal](#). The first two clustering are available in package [factanal](#).
- Package [factanal](#) provides computational hierarchical clustering which can specifically assign the minority data to incorrect structure present in the data set from 'weak' genes.
- Package [factanal](#) provides a fast implementation of the 2003-AS density hierarchical clustering of applications with novel algorithm using a kd-tree.

If you have **VERY** large datasets

Computing the distance between each pair of datapoint is NOT reasonable

Function `clara()` is a wrapper to `pam()` for larger data sets.


Partitioning method (k-medoids)

It uses **subsamples** of the original dataset,
looks for **conserved clusters**,
Finds their **medoids**,
Compute the **distance of the remaining data points** to the medoids,
Attribute a cluster to each datapoint.

Goals for this lecture

1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as *k-means* or *k-medoids*
(also mentioned density based clustering: however cursed)
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called *hierarchical clustering*.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.

Summary: the clustering workflow



1. Start with the data.
What are the **measurements**?
What **type** of data?
2. Define or select a **metric (distance)** to evaluate the (dis)similarity between 2 samples.
3. Choose a **clustering method**:
bottom-up (*hierarchical*) or
top-down (*k-methods*)
4. **Validate** the clustering and
evaluate the **optimal number** of
clusters
5. **Augment** your data with the
cluster information

Distances

We saw at the start of the Lecture how finding the **right** distance is an essential first step in a clustering analysis. This is a case when the ***garbage in, garbage out*** motto is in full force. Always choose a distance which is **scientifically meaningful** and compare output from as many distances as possible. Sometimes the same data require **different distances** when **different scientific objectives** are pursued.

Partitioning and Aggregating

We saw two different types of clustering approaches:

iterative partitioning approaches such as kmeans and kmedoids (PAM) that alternated between estimating the clusters and assigning points to them and **hierarchical clustering** approaches that agglomerate points and then small clusters into larger ones in a nested sequence of sets that can be represented by hierarchical clustering trees.

Cluster validation

Clustering algorithms ***always*** deliver clusters so we need to assess their quality and the number of clusters to choose carefully.

These validation steps are done using visualization tools and repeating the clustering on many resamples of the data. We saw how statistics such as the bss/wss or $\log(\text{wss})$ can be calibrated using simulation on data where we understand the group structure and can provides useful benchmarks for choosing the number of cluster.

The choice or definition of distance depends on the data

Euclidean distance (L2)

$$d(A, B) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots + (a_p - b_p)^2}.$$

Manhattan distance (L1)

$$d(A, B) = |a_1 - b_1| + |a_2 - b_2| + \dots + |a_p - b_p|.$$

Maximum distance (L ∞)

$$d_\infty(A, B) = \max_i |a_i - b_i|.$$

Minkowski distance (L m)

$$d(A, B) = ((a_1 - b_1)^m + (a_2 - b_2)^m + \dots + (a_p - b_p)^m)^{\frac{1}{m}}.$$

Edit (Hamming) distance

Binary distance

Jaccard distance

$$d_J(S, T) = 1 - J(S, T) = \frac{f_{01} + f_{10}}{f_{01} + f_{10} + f_{11}}.$$






Correlation-based distance

$$d(A, B) = \sqrt{2(1 - \text{cor}(A, B))}.$$

Weighted Euclidean distance

Mahalanobis distance

...



Questions

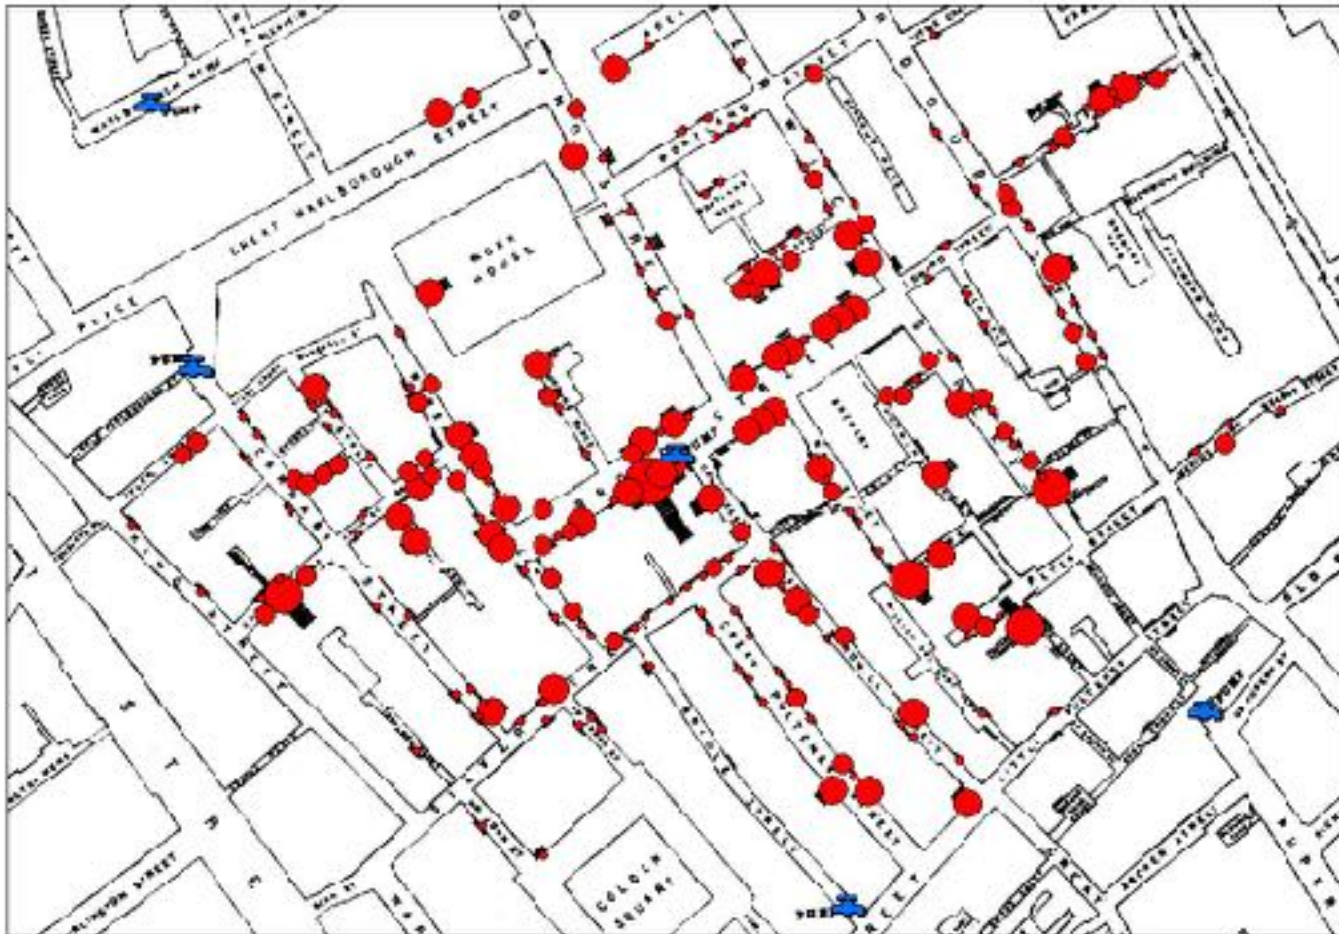


Goals for this lecture



1. Review measures of **(dis)similarity** that help us define clusters.
2. Explain **non-parametric methods** such as ***k-means*** or ***k-medoids***
3. Explain the **recursive approach** to clustering that combines observations and groups into a hierarchy of sets and called ***hierarchical clustering***.
4. Understand how to **validate the clusters** and select the **optimal number of clusters**.
5. Examples in R - lab is important.
6. Why do we cluster ?

Why do we cluster ?

- Hidden variables:



Why do we cluster ?



Why do we cluster ?

- Hidden variables:



